

VM5k and DVMS on Grid'5000

Deploying and Managing Thousands of Virtual Machines on
Hundreds of Nodes Distributed Geographically

Javier Pastor¹ Laurent Pouilloux²

¹Héméra Phd
ASCOLA - Mines Nantes / Inria

²Hemera Engineer
Inria / ENS Lyon

18-06-2014 / Grid'5000 School

Context

Cloud computing has become very popular.

- Ever-increasing demand \Rightarrow ever-increasing infrastructure size.
- PB: scalability, reliability, energy but also security, jurisdiction and network overhead.
- Decentralise the production of computing ressources
(Discovery project, <http://beyondtheclouds.github.io/>)


Scientific question

Discovery: leverage the concept of micro/nano datacenters [Greenberg2009] geographically spread. \Rightarrow nodes can be far from each other.

- And we want to maximise cooperation between close nodes/micro DCs.
- Example: The DVMS case.

Grid'5000 as a testbed

- renater interconnected network
- various hardware
- KaVLAN
- experimentfull stack control



Experimental Workflow

- 1 reserve many nodes on different sites, with a global-KaVLAN
- 2 deploy thousands of Virtual Machines
- 3 initiate stress process on them
- 4 install DVMS
- 5 use vivaldi to compute hosts distances
- 6 generate random stress on the virtual machiens
- 7 collect results

Automatic time slot selection

Virtual Machines deployment

Stress initialization

Vivaldi

Live visualization

Results Analysis

Conclusion

* publi Europar * github repository